KINGDOM AI OS

Deep Tab Guide + Pirate Face Uncensored Integration + Interface Spec

Addendum — 2026-09-30
Super in-depth every tab • Sub-tabs • LLMs • Containers • Security • Pirate Face

1. Pirate Face (pirateface.co/models) — Uncensored / Obliterated Priority

Pirate Face is a high-volume open-model index (669k+ entries) that prioritizes magnets, Hugging Face mirrors, and especially community uncensored / abliterated / heretic / decensored variants. Searches for “uncensored”, “abliterated”, “heretic” surface many Qwen3.x, GLM-5.3, DeepSeek V4.x, and other fine-tunes with dramatically reduced refusal rates. This is now a primary discovery source for the Kingdom model registry.

Integration Rule for Coders / Manus / Codex

- When adding or recommending any base model (Qwen, GLM, DeepSeek, etc.), always check Pirate Face first for an Uncensored / Heretic / Abliterated / Obliterated GGUF or safetensors variant.
- Prefer those variants for Research, Legal, OSINT, and Sovereignty workspaces.
- Record the exact Pirate Face / HF / magnet link in models.yaml with guardrail_class: uncensored or low.
- Colibrì and llama.cpp paths should point to the uncensored weights when available and license-compatible.
- Example high-priority families already appearing: Qwen3.8-27B-Uncensored / Heretic variants (HauhauCS, DavidAU, 0bserverx, etc.), GLM-5.3-Flash-Abliterated / Heretic, DeepSeek-V4.1-Flash-Abliterated, OrcaSAQ Cyber Uncensored, Ternary-Bonsai Uncensored-Heretic.

Status: Added to master inventory as primary uncensored discovery source. Router will prefer these when the workspace requests low-guardrail / uncensored.

2. Codex Integration (OpenAI Codex CLI) + Kaggle / Cloudflare / Grok

Codex CLI is OpenAI's terminal coding agent (open-source client, models behind ChatGPT plans). It can read/edit code, run commands, and hand off to cloud tasks. We treat it as one of the coding agents inside the Coding / Development workspace, alongside OpenHands, Cline, Aider, and custom Kingdom agents.

Attachment Plan

- Local install of Codex CLI inside the Coding sandbox / disposable workspace.
- Kaggle: use as GPU worker for heavy model inference, fine-tuning experiments, or large context jobs that Codex or other agents launch.
- Cloudflare: front-door for the entire Kingdom UI and API gateway (DNS, WAF, Tunnel). Codex itself stays local or in the control plane.
- Grok: available via OpenRouter or direct xAI endpoint as one of the models the router and coding agents can call. Grok can also be used as a high-reasoning model for planning steps that Codex then executes.
- All of the above is orchestrated by the Kingdom model router + agent manager so Codex does not become a single point of failure.

3. Interface Diagram Description (Based on Uploaded UI)

Dark theme. Left sidebar (collapsible): model selector (top) → Search → Home → Chat (global) → Prompt Library → Integrations → Pinned (color-coded by workspace) → History. Main area: Welcome cards or active workspace content. Top or secondary bar: the 24 workspace tabs (or grouped). Each workspace has its own full chatbot interface + Agents + Tasks + Knowledge + Tools + Files + Settings. Bright accent colors (orange, teal, pink, purple, green) for tabs and status chips. White content text. Global overlays for Active Tasks and Security alerts.

Every tab has its own chatbot: Yes. Clicking a tab switches the entire context (sandbox, tools, preferred models, knowledge base) while the global Home chat remains available. Tasks are global but filterable by workspace color.

4. Super In-Depth Every Tab (Structure + Sub-Areas + LLMs + Container + Security)

Common internal layout for every workspace: Chat (workspace-specific) | Agents | Active Tasks | Closed Tasks | Knowledge / RAG | Files / Uploads | Tools | Models (preferred + global) | Outputs | Automations | Settings. Each workspace runs in its own container/sandbox with limited mounts, network policy, and permission profile. Model router still prefers local uncensored (Pirate Face variants) when the tab requests it.

1. ■ Home / Command Center

Global chat, system dashboard, agent status, notifications, quick-launch to any workspace, overall task queue. Container: control-plane only. Security: highest — no external tools without broker.

2. ■ LLM / Models

Full registry (including all Pirate Face uncensored/heretic/abliterated variants), health, benchmarks, Colibrì / GGUF / OpenRouter pools, routing rules, evaluation lab. Sub: Local / API / Free / Uncensored / Vision / Coding. Container: model-serving isolation.

3. ■ Research

Deep research agents, academic/web/news, MinerU + PaddleOCR documents, SearXNG, Agent Reach, Crawl4AI/Firecrawl, citations. Preferred LLMs: uncensored reasoning MoEs (GLM/DeepSeek/Qwen heretic variants via Colibrì or GGUF). Container: research sandbox with outbound web policy.

4. ■■ Legal

Cases, claims, litigation, sovereignty/constitutional, treaties, evidence library (access-controlled), document intelligence. Preferred: high-reasoning uncensored models. Container: legal sandbox + encrypted evidence store. Strict audit.

5. ■■ OSINT

Investigation workflows, people/companies, threat intel, exposure monitoring, Agent Reach + Browser Use + SearXNG + specialized sources (OpenCorporates, OpenAlex, ICIJ, Old Maps Online, etc.). Container: OSINT sandbox with careful network and credential broker.

6–7. ■ Finance / ■ Investments

Analysis, modeling, portfolio (interfaces only). No direct ledger. Preferred: reasoning + code models. Container: finance interface sandbox. All real money moves go through separate Financial Platform API with multi-approval.

8. ■ Real Estate

Properties, research, maps, valuation, documents. Tools: scraping, maps, OCR. Container: real-estate sandbox.

9–10. ■ Business / CRM + ■ Leads

Contacts, pipelines, enrichment, outreach (Listmonk etc.). Container: CRM sandbox. Outreach agents require approval for high-volume or external sends.

11–12. ■■ City Planning + ■■ Engineering

Master planning, CAD/tech, robotics, simulation. Container: engineering sandbox with possible GPU mount.

13. ■ Coding / Development

OpenHands, Codex CLI, Cline, Aider, custom agents, disposable workspaces, app factory, Dify. Preferred coding models (DeepSeek, Qwen-Coder, local uncensored). Container: per-agent disposable sandbox. Kaggle attachable for heavy jobs.

14. ■ Web / Scraping

Browser Use (primary full browser), Agent Reach (zero-key platform access), Playwright, Crawl4AI, Firecrawl, SearXNG. Credential Broker mandatory for any authenticated session. Container: web sandbox with network policy.

15–16. ■ News + ■ Social / Marketing

RSSHub, Folo, aggregation, publishing, campaigns. Container: media/comms sandbox.

17–18. ■ Fashion + ■ Images / Visual AI

ComfyUI, FLUX + IP-Adapter, InvokeAI, Human Studio (Character ID + diverse skin/body/hair consistency + verification). Preferred vision models. Container: media GPU sandbox.

19. ■ Video / UGC / TV / Movies

VACE (Wan2.1-VACE), Wan, LTX, Open-Sora, ComfyUI, remake engine (shot/pose/depth/camera → target character), FFmpeg, talking-head systems. Container: video GPU sandbox (Kaggle or local).

20. ■ Music / Audio

Stem tools, Whisper.cpp, F5-TTS / GPT-SoVITS, generation, mixing. Container: audio sandbox.

21. ■ Agents / Communications

Agent Creator + library (100+ conceptual agents), identities/faces, phone/email/messaging, Agent Reach + Browser Use skills. Container: agent runtime with strict permission profiles.

22. ■ Tasks

Global open / in-progress / completed / failed / scheduled. Color-coded by originating workspace. Click jumps to the tab. Container: control-plane.

23. ■ Open Source / Apps

Full tool registry (including Pirate Face, Agent Reach, Browser Use, MinerU, VACE, etc.), Dify, MCP, plugins, app factory. Container: registry service.

24. ■ Security / Admin

Keycloak, secrets (OpenBao), Wazuh, CrowdSec, Suricata, audit logs, threat profiling, backups, sandbox management. Container: security plane (highest isolation). Continuous self-monitoring of security tools.

5. Container & Never-Break Security Summary

- Each major workspace (or group) runs in its own container with least-privilege mounts, network policy, and resource limits.
- Model serving (Colibrì, llama.cpp, vLLM) is further isolated; agents call via the model router, not direct process access.
- Credential Broker: agents request → human or policy approval for high-impact → short-lived scoped secret → action → receipt → audit.
- Financial Platform: completely outside. Only controlled, audited APIs.
- Plugin lifecycle: register → scan (incl. Pirate Face source) → sandbox test → benchmark → permission profile → deploy → health → auto-rollback on failure.
- Pirate Face uncensored models are treated as first-class but still go through the same scan + evaluation before production routing.

End of Deep Tab Guide + Pirate Face Addendum

Use together with the Master Architecture Bible PDF and the starter ZIP package.